

Other properties

- Every single variable has a univariate normal distribution.

$$x_j \sim \mathcal{N}(\mu_j, \sigma_{jj}^2)$$

$$\begin{bmatrix} x_1 \\ \vdots \\ x_p \end{bmatrix} \sim \mathcal{N}(\mu, \Sigma)$$

- Any subset of the variables also has a multivariate normal distribution.

$$\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \sim \mathcal{N} \left(\begin{bmatrix} \mu_1 \\ \mu_2 \\ \mu_2 \end{bmatrix}; \begin{bmatrix} \end{bmatrix} \right)$$

- Zero covariance terms or a diagonal covariance matrix implies that the variables are independent of each other.

$$\Sigma_{3 \times 3} = \begin{bmatrix} \sigma_{11} & 0 & 0 \\ 0 & \sigma_{22} & 0 \\ 0 & 0 & \sigma_{33} \end{bmatrix}$$

$$x_1 \perp\!\!\!\perp x_2 \quad x_3 \perp\!\!\!\perp x_1 \\ x_2 \perp\!\!\!\perp x_3$$

- Any conditional distribution for a subset of the variables conditional on known values for another subset of variables is a multivariate distribution.

Partitioned Gaussian Distributions

$$p(\underline{\mathbf{x}}) = \mathcal{N}(\underline{\mathbf{x}} | \underline{\boldsymbol{\mu}}, \underline{\boldsymbol{\Sigma}})$$

$\leftarrow P_{X1} \leftarrow P_{1 \times P}$

$$X_a \cap X_b = \emptyset$$

$$\underline{\mathbf{x}} = \begin{pmatrix} \underline{\mathbf{x}}_a \\ \underline{\mathbf{x}}_b \end{pmatrix}$$

$$\underline{\boldsymbol{\mu}} = \begin{pmatrix} \underline{\boldsymbol{\mu}}_a \\ \underline{\boldsymbol{\mu}}_b \end{pmatrix}$$

$$\underline{\boldsymbol{\Sigma}} = \begin{pmatrix} \underline{\boldsymbol{\Sigma}}_{aa} & \underline{\boldsymbol{\Sigma}}_{ab} \\ \underline{\boldsymbol{\Sigma}}_{ba} & \underline{\boldsymbol{\Sigma}}_{bb} \end{pmatrix}$$

$$\underline{\boldsymbol{\Lambda}} \equiv \underline{\boldsymbol{\Sigma}}^{-1}$$

$$\underline{\boldsymbol{\Lambda}} = \begin{pmatrix} \boldsymbol{\Lambda}_{aa} & \boldsymbol{\Lambda}_{ab} \\ \boldsymbol{\Lambda}_{ba} & \boldsymbol{\Lambda}_{bb} \end{pmatrix}$$

$$X_a \sim \mathcal{N}(\mu_a; \Sigma_{aa}) :$$

Partitioned Conditionals and Marginals

$$p(\mathbf{x}_a | \mathbf{x}_b) = \mathcal{N}(\mathbf{x}_a | \underline{\mu}_{a|b}, \underline{\Sigma}_{a|b})$$

$$\underline{\Sigma}_{a|b} = \underline{\Lambda}_{aa}^{-1} = \underline{\Sigma}_{aa} - \underline{\Sigma}_{ab} \underline{\Sigma}_{bb}^{-1} \underline{\Sigma}_{ba}$$

$$\underline{\mu}_{a|b} = \underline{\Sigma}_{a|b} \{ \underline{\Lambda}_{aa} \underline{\mu}_a - \underline{\Lambda}_{ab} (\mathbf{x}_b - \underline{\mu}_b) \}$$

$$= \underline{\mu}_a - \underline{\Lambda}_{aa}^{-1} \underline{\Lambda}_{ab} (\mathbf{x}_b - \underline{\mu}_b)$$

$$= \underline{\mu}_a + \underline{\Sigma}_{ab} \underline{\Sigma}_{bb}^{-1} (\mathbf{x}_b - \underline{\mu}_b)$$

$$p(\mathbf{x}_a) = \int p(\mathbf{x}_a, \mathbf{x}_b) d\mathbf{x}_b$$

$$= \mathcal{N}(\mathbf{x}_a | \underline{\mu}_a, \underline{\Sigma}_{aa})$$

$$a = \{1\}; \quad b = \{2\}$$

$$P(x_1 | x_2) = \mathcal{N}(\mu_{1|2}; \Sigma_{1|2})$$

$$P\left(\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}\right)$$

Derive for bi-variate case.

$$\underline{\Sigma} = \begin{bmatrix} \sigma_{11}^2 & \sigma_{12} \\ \sigma_{21} & \sigma_{22}^2 \end{bmatrix}$$


$$\underline{\mu} = \begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix}$$

$$\sigma_{1|2}^2 = \sigma_{11}^2 - \frac{\sigma_{12}^2}{\sigma_{22}^2} \quad \checkmark$$


$$\mu_{1|2} = \mu_1 + \frac{\sigma_{12}}{\sigma_{22}^2} (x_2 - \mu_2)$$

Conditional distribution for bivariate case

$$\text{Mean} = \mu_1 + \frac{\sigma_{12}}{\sigma_{22}}(x_2 - \mu_2)$$

$$\text{Variance} = \sigma_{11} - \frac{\sigma_{12}^2}{\sigma_{22}}$$
A hand-drawn red bracket is positioned below the variance formula, spanning from the σ_{11} term to the end of the fraction. Below the bracket is a horizontal red line.

Partitioned Conditionals and Marginals



Demos: <https://colab.research.google.com/github/goodboychan/goodboychan.github.io/blob/main/notebooks/2021-08-11-Multivariate-distribution.ipynb>

Example 6-1: Conditional Distribution of Weight Given Height for College Men

Suppose that the weights (lbs) and heights (inches) of undergraduate college men have a multivariate normal distribution with mean vector $\mu = \begin{pmatrix} 175 \\ 71 \end{pmatrix}$ and covariance matrix $\Sigma = \begin{pmatrix} 550 & 40 \\ 40 & 8 \end{pmatrix}$.


$$\mu_{kg} = \begin{pmatrix} 80 \\ 71 \end{pmatrix}$$

$$\Sigma_{kg} = \begin{pmatrix} 550/4 & 20 \\ 20 & 8 \end{pmatrix}$$

The conditional distribution of X_1 weight given $x_2 = \text{height}$ is a normal distribution with

$$\begin{aligned} \text{Mean} &= \mu_1 + \frac{\sigma_{12}}{\sigma_{22}}(x_2 - \mu_2) \\ &= 175 + \frac{40}{8}(x_2 - 71) \\ &= -180 + 5x_2 \end{aligned}$$

$$\begin{aligned} \text{Variance} &= \sigma_{11} - \frac{\sigma_{12}^2}{\sigma_{22}} \\ &= 550 - \frac{40^2}{8} \\ &= 350 \end{aligned}$$



For instance, for men with height = 70, weights are normally distributed with mean = $-180 + 5(70) = 170$ pounds and variance = 350. (So standard deviation $\sqrt{350} = 18.71 = \text{pounds}$)

Notice that we have generated a simple linear regression model that relates weight to height.

Geometry of the Multivariate Normal Distribution

- Can we characterize the shape and orientation of the ellipse that defines that contours of equal density?

Constant probability density contour = {all \mathbf{x} such that $(\mathbf{x} - \boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) = c^2$ }




Figure 4.4 The 50% and 90% contours for the bivariate normal distributions in Figure 4.2.

- We will see that these can be characterized using eigen vectors and values of the covariance matrix.

Eigen values and Eigen vectors

- A square matrix A has a eigen value, eigen vector pair λ , $e \neq 0$ if $Ae = \lambda e$ where norm of e is 1

Let **A** be a $k \times k$ square symmetric matrix. Then **A** has k pairs of eigenvalues and eigenvectors namely,

$$\lambda_1, e_1 \quad \lambda_2, e_2 \quad \dots \quad \lambda_k, e_k \quad (2-15)$$

The eigenvectors can be chosen to satisfy 1 = $e_1' e_1$ = \dots = $e_k' e_k$ and be mutually perpendicular. The eigenvectors are unique unless two or more eigenvalues are equal.

$$e_i^T e_i = e_i' e_i = \langle e_i, e_i \rangle = \sum_{j=1}^p e_{ij}^2$$
$$e_j^T e_k = 0 \quad \forall j \neq k$$

Spectral decomposition of A

$$\underset{(k \times k)}{\mathbf{A}} = \lambda_1 \underset{(k \times 1)}{e_1} \underset{(1 \times k)}{e_1'} + \lambda_2 \underset{(k \times 1)}{e_2} \underset{(1 \times k)}{e_2'} + \dots + \lambda_k \underset{(k \times 1)}{e_k} \underset{(1 \times k)}{e_k'}$$

Spectral decomposition of a positive semi-definite matrix

- If A is positive-definite then all eigen-values ≥ 0

- Example:

$$\mathbf{R} = \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix}$$

$y^T A y \geq 0$
choose $y = e_j$ to show
that $\lambda_j \geq 0$

- First find Eigen values and vectors.

- [4.5 - Eigenvalues and Eigenvectors | STAT 505 \(psu.edu\)](#)

[HW]

$$e_1 = \begin{pmatrix} \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \end{pmatrix} \text{ for } \lambda_1 = 1 + \rho \text{ and } \begin{pmatrix} \frac{1}{\sqrt{2}} \\ -\frac{1}{\sqrt{2}} \end{pmatrix} \text{ for } \lambda_2 = 1 - \rho$$

$e_1 \cdot e_2 = 0$

$$\begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix} = (1 + \rho) \begin{pmatrix} \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \end{pmatrix} \begin{pmatrix} \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \end{pmatrix} + (1 - \rho) \begin{pmatrix} \frac{1}{\sqrt{2}} \\ -\frac{1}{\sqrt{2}} \end{pmatrix} \begin{pmatrix} \frac{1}{\sqrt{2}} & -\frac{1}{\sqrt{2}} \end{pmatrix}$$

$$\text{If } \Sigma = \lambda_1 e_1 e_1^T + \dots + \lambda_p e_p e_p^T$$

Then

$$\Sigma^{-1} = \frac{1}{\lambda_1} e_1 e_1^T + \dots + \frac{1}{\lambda_p} e_p e_p^T$$

Geometry of the Multivariate Gaussian

$$\underline{C^2} = \underline{\Delta^2} = (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})$$


$$\underline{\boldsymbol{\Sigma}^{-1}} = \sum_{i=1}^p \frac{1}{\lambda_i} \mathbf{u}_i \mathbf{u}_i^T \quad \sum_{i=1}^p \frac{1}{\lambda_i} \mathbf{e}_i \mathbf{e}_i^T$$

$$C^2 = \Delta^2 = \sum_{i=1}^p \frac{y_i^2}{\lambda_i}$$

$$y_i = \mathbf{u}_i^T (\mathbf{x} - \boldsymbol{\mu}) \quad \mathbf{e}_i^T (\mathbf{x} - \boldsymbol{\mu})$$

$$(\mathbf{x} - \boldsymbol{\mu})^T \left[\sum_{i=1}^p \frac{1}{\lambda_i} \mathbf{e}_i \mathbf{e}_i^T \right] (\mathbf{x} - \boldsymbol{\mu})$$

$$\Delta^2 = \sum_{i=1}^p \frac{1}{\lambda_i} (\mathbf{e}_i^T (\mathbf{x} - \boldsymbol{\mu}))^2$$




Principal component analysis

Projecting high-dimensional data

- When multivariate dataset has a large number of variables, analysis and interpretation of the data may be hard.
- Too many variables pairs, so pairwise correlation may be hard to grasp.
- For convenient visualization and interpretation
 - Reduce the number of variables.
- How to reduce number of variables (while capturing most of the information in the data)
 - Information == variance

Example



What is the best way to summarize this two dimensional data into a single dimension without losing much of the dispersion?

How to reduce number of variables: many methods

- Principal component analysis ~~←~~
- Factor analysis
- Other embedding methods
 - Random projection
 - T-SNE ~~△~~

Principal component analysis

- Let original set of p variables be X_1, X_2, \dots, X_p
- Define a smaller set of new variables that are linear combinations of existing variables.

$$\begin{array}{rcl} Y_1 & = & e_{11}X_1 + e_{12}X_2 + \dots + e_{1p}X_p \\ Y_2 & = & e_{21}X_1 + e_{22}X_2 + \dots + e_{2p}X_p \\ & \vdots & \\ Y_p & = & e_{p1}X_1 + e_{p2}X_2 + \dots + e_{pp}X_p \end{array}$$

Variance and Co-variance of the new variables.

Let

$$\text{var}(\underline{\mathbf{X}}) = \Sigma = \begin{pmatrix} \sigma_1^2 & \sigma_{12} & \dots & \sigma_{1p} \\ \sigma_{21} & \sigma_2^2 & \dots & \sigma_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_{p1} & \sigma_{p2} & \dots & \sigma_p^2 \end{pmatrix}$$

Then:

$$\text{var}(\underline{Y_i}) = \sum_{k=1}^p \sum_{l=1}^p e_{ik} e_{il} \sigma_{kl} = \underline{\mathbf{e}_i}' \Sigma \underline{\mathbf{e}_i}$$

$$Y_i = \begin{bmatrix} e_{i1} \\ \vdots \\ e_{ip} \end{bmatrix}^T \begin{bmatrix} x_1 \\ \vdots \\ x_p \end{bmatrix}$$

$$\text{cov}(\underline{Y_i}, \underline{Y_j}) = \sum_{k=1}^p \sum_{l=1}^p e_{ik} e_{jl} \sigma_{kl} = \underline{\mathbf{e}_i}' \Sigma \underline{\mathbf{e}_j}$$

Principal components

- First principal component Y_1 is chosen to maximize the variance among all possible linear combinations such that the norm of coefficients is 1.

More formally, select $\mathbf{e}_{11}, \mathbf{e}_{12}, \dots, \mathbf{e}_{1p}$ that maximizes

$$\text{var}(Y_1) = \sum_{k=1}^p \sum_{l=1}^p e_{1k} e_{1l} \sigma_{kl} = \mathbf{e}_1' \Sigma \mathbf{e}_1$$

subject to the constraint that

$$\mathbf{e}_1' \mathbf{e}_1 = \sum_{j=1}^p e_{1j}^2 = 1$$

Second principal component

Select $\underline{e_{21}}, \underline{e_{22}}, \dots, \underline{e_{2p}}$ that maximizes the variance of this new component...

$$\text{var}(\underline{Y_2}) = \sum_{k=1}^p \sum_{l=1}^p e_{2k} e_{2l} \sigma_{kl} = \underline{e_2}' \Sigma \underline{e_2}$$

subject to the constraint that the sums of squared coefficients add up to one,

$$\underline{e_2}' \underline{e_2} = \sum_{j=1}^p e_{2j}^2 = 1$$

along with the additional constraint that these two components are uncorrelated.

$$\text{cov}(\underline{Y_1}, \underline{Y_2}) = \sum_{k=1}^p \sum_{l=1}^p e_{1k} e_{2l} \sigma_{kl} = \underline{e_1}' \Sigma \underline{e_2} = 0$$

i^{th} Principal Component (PC*i*): Y_i

We select $e_{i1}, e_{i2}, \dots, e_{ip}$ to maximize

$$\text{var}(Y_i) = \sum_{k=1}^p \sum_{l=1}^p e_{ik} e_{il} \sigma_{kl} = \mathbf{e}_i' \Sigma \mathbf{e}_i$$

subject to the constraint that the sums of squared coefficients add up to one...along with the additional constraint that this new component is uncorrelated with all the previously defined components.

$$\rightarrow \mathbf{e}_i' \mathbf{e}_i = \sum_{j=1}^p e_{ij}^2 = 1$$

$$\text{cov}(Y_1, Y_i) = \sum_{k=1}^p \sum_{l=1}^p e_{1k} e_{il} \sigma_{kl} = \mathbf{e}_1' \Sigma \mathbf{e}_i = 0,$$

$$\text{cov}(Y_2, Y_i) = \sum_{k=1}^p \sum_{l=1}^p e_{2k} e_{il} \sigma_{kl} = \mathbf{e}_2' \Sigma \mathbf{e}_i = 0,$$

\vdots

$$\text{cov}(Y_{i-1}, Y_i) = \sum_{k=1}^p \sum_{l=1}^p e_{i-1,k} e_{il} \sigma_{kl} = \mathbf{e}_{i-1}' \Sigma \mathbf{e}_i = 0$$

For what Y_1 is $\text{Variance}(Y_1)$ maximized?

- The coefficient of the first principal component correspond to the Eigen vector with the maximum Eigen value.

More generally

- The i -th principal component corresponds the i -th largest eigen vector.

The variance for the i th principal component is equal to the i th eigenvalue.

$$\text{var}(Y_i) = \text{var}(e_{i1}X_1 + e_{i2}X_2 + \dots e_{ip}X_p) = \lambda_i$$

$$\text{cov}(Y_i, Y_j) = 0$$

The proportion of variance explained

- The total variance of X
- We can show that sum of p Eigen values equals the total variance

- The fraction of variance explained by the i -th Eigen value $\frac{\lambda_i}{\lambda_1 + \lambda_2 + \cdots + \lambda_p}$

Reducing number of dimensions

- Variance explained by first k Eigen values $\frac{\lambda_1 + \lambda_2 + \dots + \lambda_k}{\lambda_1 + \lambda_2 + \dots + \lambda_p}$

11.3 - Example: Places Rated

Example 11-2: Places Rated

We will use the Places Rated Almanac data (Boyer and Savageau) which rates 329 communities according to nine criteria:

1. Climate and Terrain
2. Housing
3. Health Care & Environment
4. Crime
5. Transportation
6. Education
7. The Arts
8. Recreation
9. Economics

[11.3 - Example: Places Rated | STAT 505 \(psu.edu\)](#)

Notes

- The data for many of the variables are strongly skewed to the right.
- The log transformation was used to normalize the data.

More demos

- <https://colab.research.google.com/github/jakevdp/PythonDataScienceHandbook/blob/master/notebooks/05.09-Principal-Component-Analysis.ipynb>